

ggfunction: A Grammar of Graphics for Mathematical Functions and Probability Distributions

by Dusty Turner, David Kahle, Rodney X. Sturdivant, and James Otto

Abstract The `ggfunction` R package extends `ggplot2` by providing a principled, unified interface for visualizing mathematical functions, probability distributions, and empirical data distributions. It is organized around three complementary families. The first is a dimensional taxonomy classifying functions by their input and output dimensions: scalar functions $f: \mathbb{R} \rightarrow \mathbb{R}$, parametric curves $\gamma: \mathbb{R} \rightarrow \mathbb{R}^2$, scalar fields $f: \mathbb{R}^2 \rightarrow \mathbb{R}$, and vector fields $\mathbf{F}: \mathbb{R}^2 \rightarrow \mathbb{R}^2$. The second family provides specialized geoms for theoretical probability distributions—density, cumulative distribution, mass, quantile, discrete cumulative distribution, discrete quantile, discrete survival, survival, hazard, and cumulative hazard functions, together with bivariate density and mass functions—each with built-in support for region shading central to hypothesis testing and interval estimation. The third family provides geoms for empirical distributions computed from samples: the empirical CDF with a simultaneous Kolmogorov–Smirnov confidence band, the empirical quantile function with the inverted KS band, one-sample PP, QQ, and stabilized probability (SP) diagnostic plots with DKW bands under a fully specified null distribution, the empirical PMF as a lollipop chart, and the empirical cumulative hazard function with a transformed DKW confidence band. For right-censored survival data, `geom_ecdf_km()` and `geom_echf_na()` provide the Kaplan–Meier survival estimator (Kaplan and Meier, 1958) and Nelson–Aalen cumulative hazard estimator (Nelson, 1972; Aalen, 1978), respectively, each with simultaneous equal-precision confidence bands (Nair, 1984). By embedding function evaluation, numerical integration, and validation directly into the `ggplot2` layer system, `ggfunction` lets users move from a mathematical definition to a publication-quality visualization in a single, composable call.

1 Introduction

Visualizing mathematical functions is one of the most common tasks in applied mathematics and statistics—and one of the most under-supported in the grammar of graphics. Whether an instructor is shading a rejection region under a normal curve, a researcher is inspecting a scalar field over a spatial domain, or a student is tracing a parametric spiral, the need to move quickly from a function’s definition to its graphical representation arises constantly. In R, `ggplot2` (Wickham, 2010, 2016) has become the dominant visualization framework, yet its native support for plotting functions remains narrow. The built-in `stat_function()` evaluates a univariate function over a range and draws it as a path—perfectly adequate for $f: \mathbb{R} \rightarrow \mathbb{R}$, but offering nothing for parametric curves, scalar fields, vector fields, or the rich family of functions associated with probability distributions.

Several packages address pieces of this gap. `ggdist` (Kay, 2024) provides sophisticated distribution visualizations oriented toward Bayesian uncertainty communication. `mosaic` (Pruim et al., 2017) offers `plotDist()` for pedagogical work, but it is built on the `lattice` (Sarkar, 2008) framework rather than the grammar of graphics. `metR` provides contour and vector-field displays for meteorological data, though it requires precomputed grids rather than function objects. None of these packages offers a unified treatment of functions organized by their dimensional structure.

`ggfunction` fills this gap by extending `ggplot2` with a principled taxonomy of function types, classified by the dimensions of their domain and codomain. It provides three complementary families of geoms:

1. A **dimensional taxonomy** covering four function types: $\mathbb{R} \rightarrow \mathbb{R}$, $\mathbb{R} \rightarrow \mathbb{R}^2$, $\mathbb{R}^2 \rightarrow \mathbb{R}$, and $\mathbb{R}^2 \rightarrow \mathbb{R}^2$.
2. A **probability distribution family** providing specialized layers for density, cumulative distribution, mass, quantile, discrete cumulative distribution, discrete quantile, discrete survival, survival, hazard, and cumulative hazard functions, as well as bivariate density and mass functions visualized through highest density regions and lattice displays.
3. An **empirical data family** providing layers for the empirical CDF, empirical quantile function, one-sample PP, QQ, and stabilized probability (SP) diagnostic plots, empirical PMF, and empirical cumulative hazard function, each computed directly from a sample and accompanied by simultaneous DKW/Massart confidence bands where appropriate. For right-censored survival data, the family additionally provides the Kaplan–Meier survival estimator (Kaplan and Meier, 1958) with a simultaneous equal-precision Greenwood band (Greenwood, 1926; Nair, 1984).

and the Nelson–Aalen cumulative hazard estimator (Nelson, 1972; Aalen, 1978) with pointwise normal bands.

Each layer follows the **ggplot2** extension pattern: a Stat object evaluates the user-supplied function on a grid or sequence (or tabulates a sample), and a corresponding Geom renders the result. Because every **ggfunction** layer is a standard **ggplot2** layer, it composes freely with the full ecosystem of scales, coordinates, themes, and facets.

The remainder of this article is organized as follows. We describe the design principles underlying the package in Section 2. Sections 3 and 4 present the dimensional taxonomy and probability distribution family, respectively, illustrating each with examples. Section 5 introduces the empirical data family. Section 6 addresses implementation details, Section 7 demonstrates composability with the broader grammar, Section 8 surveys related packages, and Section 9 concludes.

2 Design principles

2.1 A taxonomy grounded in dimension

The central organizing principle of **ggfunction** is that any function visualizable in two-dimensional space belongs to one of four classes, determined by the dimensions of its domain m and codomain n :

$$\begin{aligned} f: \mathbb{R} &\rightarrow \mathbb{R} & (m = 1, n = 1) \\ \gamma: \mathbb{R} &\rightarrow \mathbb{R}^2 & (m = 1, n = 2) \\ f: \mathbb{R}^2 &\rightarrow \mathbb{R} & (m = 2, n = 1) \\ \mathbf{F}: \mathbb{R}^2 &\rightarrow \mathbb{R}^2 & (m = 2, n = 2) \end{aligned} \tag{1}$$

Each class admits a natural visual encoding:

- **Scalar functions** ($\mathbb{R} \rightarrow \mathbb{R}$): curves in the Cartesian plane, optionally with region shading between the curve and the x -axis.
- **Parametric curves** ($\mathbb{R} \rightarrow \mathbb{R}^2$): directed paths with color encoding the time parameter.
- **Scalar fields** ($\mathbb{R}^2 \rightarrow \mathbb{R}$): raster heatmaps, contour lines colored by level value, or filled contour regions.
- **Vector fields** ($\mathbb{R}^2 \rightarrow \mathbb{R}^2$): short arrows at grid points (default) or streamlines computed by numerical integration of the field.

The naming convention—`geom_function_1d_1d()`, `geom_function_1d_2d()`, `geom_function_2d_1d()`, and `geom_function_2d_2d()`—encodes the dimensional signature directly, making the function’s type explicit at the point of use.

2.2 Integration with the grammar of graphics

The grammar of graphics (Wilkinson, 2005) decomposes a graphic into orthogonal components: data, aesthetic mappings, geometric objects, statistical transformations, scales, coordinate systems, and facets. **ggplot2** implements this decomposition through the `ggproto` object system, which allows new Stat and Geom classes to be defined and composed with existing infrastructure.

Every **ggfunction** layer follows this pattern. The user supplies a function object via the `fun` parameter and domain bounds via `xlim` and, where appropriate, `ylim`. The Stat evaluates the function on a suitable grid and returns a data frame whose columns correspond to the aesthetic variables expected by the Geom. Because the result is a standard **ggplot2** layer, it composes with `+` alongside any other layer, scale, or theme.

A key design choice is the use of `rlang::inject()` to splice additional function arguments supplied through the `args` parameter. This allows users to parameterize functions without closures or anonymous wrappers:

```
# instead of wrapping in an anonymous function:
ggplot() + geom_pdf(fun = function(x) dnorm(x, mean = 5, sd = 2), xlim = c(0, 10))

# users can write:
ggplot() + geom_pdf(fun = dnorm, xlim = c(0, 10), args = list(mean = 5, sd = 2))
```

This pattern is consistent across all **ggfunction** layers and mirrors the `args` parameter already familiar from `ggplot2::stat_function()`.




Figure 1: The sine function over one full period.




Figure 2: The cosine function over $[0, 2\pi]$ with the interval $[0, \pi/2]$ shaded; the shaded area equals $\int_0^{\pi/2} \cos(x) dx = 1$.

3 The dimensional taxonomy

3.1 Scalar functions: $\mathbb{R} \rightarrow \mathbb{R}$

The simplest case is a univariate function rendered as a curve in the plane. `geom_function_1d_1d()` generalizes `ggplot2::stat_function()` by adding support for region shading. The function is evaluated at n equally-spaced points (default $n = 101$) over the specified `xlim`, and the resulting (x, y) pairs are drawn as a path.

```
library("ggfunction")
```

```
ggplot() +  
  geom_function_1d_1d(fun = sin, xlim = c(0, 2 * pi))
```

The `shade_from` and `shade_to` parameters fill the region between the curve and the x -axis over a specified interval $[a, b]$. Boundary values are computed by linear interpolation via `stats::approxfun()`, so the shaded region aligns precisely with the requested interval even when a or b fall between evaluation points.

```
ggplot() +  
  geom_function_1d_1d(fun = cos, xlim = c(0, 2 * pi), shade_from = 0, shade_to = pi/2)
```




Figure 3: The lemniscate of Bernoulli $\gamma(t) = (\cos t / (1 + \sin^2 t), \sin t \cos t / (1 + \sin^2 t))$ —a figure-eight curve along which the product of distances to the two foci is constant—with color encoding the parameter t .

3.2 Parametric curves: $\mathbb{R} \rightarrow \mathbb{R}^2$

A parametric curve $\gamma(t) = (x(t), y(t))$ maps a scalar parameter—typically time—to a trajectory in the plane. `geom_function_1d_2d()` evaluates the user-supplied function over the range specified by `tlim` with step size `dt`, producing columns `t`, `x`, and `y`. By default, color is mapped to `after_stat(t)` to encode progression along the curve.

```
lemniscate <- function(t) c(cos(t) / (1 + sin(t)^2), sin(t) * cos(t) / (1 + sin(t)^2))

ggplot() +
  geom_function_1d_2d(fun = lemniscate, tlim = c(0, 1.9 * pi), tail_point = TRUE)
```

The `tail_point` parameter adds a marker at the starting position, useful for indicating an initial condition. The `args` parameter supports parameterized curve families. Lissajous figures $\gamma(t) = (\sin(at + \delta), \sin(bt))$ are a natural example: the frequency ratio a/b governs the curve's topology, while a single function definition handles all cases via `args`:

```
lissajous <- function(t, a = 3, b = 2, delta = pi/2) {
  c(sin(a * t + delta), sin(b * t))
}

p1 <- ggplot() +
  geom_function_1d_2d(fun = lissajous, tlim = c(0, 1.9*pi), args = list(a = 1, b = 1)) +
  ggtitle("a = 1, b = 1")

p2 <- ggplot() +
  geom_function_1d_2d(fun = lissajous, tlim = c(0, 1.9*pi), args = list(a = 2, b = 1)) +
  ggtitle("a = 2, b = 1")

p3 <- ggplot() +
  geom_function_1d_2d(fun = lissajous, tlim = c(0, 1.9*pi), args = list(a = 3, b = 2)) +
  ggtitle("a = 3, b = 2")

library("patchwork")

(p1 | p2 | p3) + plot_layout(guides = "collect")
```

3.3 Scalar fields: $\mathbb{R}^2 \rightarrow \mathbb{R}$

A scalar field assigns a real value to each point in a two-dimensional domain. `geom_function_2d_1d()` evaluates the user-supplied function on a regular $n \times n$ grid (default $n = 50$) over $xlim \times ylim$. The




Figure 4: Lissajous figures $\gamma(t) = (\sin(at + \pi/2), \sin(bt))$ for three frequency ratios a/b . Each ratio produces a qualitatively distinct closed curve; the same function definition is reused across all three panels via the `args` parameter.




Figure 5: The function $f(x, y) = \sin(x) \cos(y)$ over $[-\pi, \pi]^2$ rendered as a raster heatmap.

function must accept a numeric vector of length two and return a scalar; internally, a helper applies it row-wise over the grid. Three visualization modes are available through the `type` argument:

- "raster" (default): a heatmap with fill mapped to `after_stat(z)`.
- "contour": iso-level curves with colour mapped to `after_stat(level)`, producing a colored legend automatically.
- "contour_filled": filled regions between contour levels rendered by `ggplot2::GeomContourFilled`.

```
f_sc <- function(u) {
  x <- u[1]; y <- u[2]
  sin(x) * cos(y)
}

ggplot() +
  geom_function_2d_1d(fun = f_sc, xlim = c(-pi, pi), ylim = c(-pi, pi))
```

The contour modes are selected simply by changing `type`:

```
p_contour <- ggplot() +
  geom_function_2d_1d(
    fun = f_sc, xlim = c(-pi, pi), ylim = c(-pi, pi), type = "contour"
  ) + ggtitle("type = 'contour'")

p_filled <- ggplot() +
  geom_function_2d_1d(
    fun = f_sc, xlim = c(-pi, pi), ylim = c(-pi, pi), type = "contour_filled"
  ) + ggtitle("type = 'contour_filled'")
```




Figure 6: The same function rendered as contour lines (left) and filled contours (right).




Figure 7: The rotation field $F(x, y) = (-y, x)$ displayed as short arrows at each grid point.

p_contour | p_filled

For the contour variants, the `StatFunction2dContour` and `StatFunction2dContourFilled` classes extend the corresponding `ggplot2` stats. The `setup_params()` method pre-computes the function grid so that `z.range` is available for automatic break computation before any data is rendered.

3.4 Vector fields: $\mathbb{R}^2 \rightarrow \mathbb{R}^2$

A vector field $F(x, y) = (F_1(x, y), F_2(x, y))$ assigns a direction and magnitude to each point in the plane. `geom_function_2d_2d()` delegates to `ggvfields` (Turner et al., 2026) and supports two display modes through the `type` argument.

Vector arrows (default). With `type = "vector"` (the default), short arrows are drawn at each grid point, oriented according to the field direction and scaled relative to the grid spacing.

```
f_rotation <- function(u) {
  x <- u[1]; y <- u[2]
  c(-y, x)
}

ggplot() +
  geom_function_2d_2d(fun = f_rotation, xlim = c(-1, 1), ylim = c(-1, 1))
```

Streamlines. With `type = "stream"`, integral curves are computed by numerical integration of the field using a fourth-order Runge-Kutta method (Turner et al., 2026) and rendered as directed paths seeded on a regular grid.

```
ggplot() +
  geom_function_2d_2d(fun = f_rotation, xlim = c(-1, 1), ylim = c(-1, 1),
    type = "stream")
```




Figure 8: The same rotation field rendered as streamlines, each following the counterclockwise flow induced by the field.

Table 1: Cross-conversion routes for distribution geoms. Each geom accepts its native function type via `fun` and one or more alternate types via the listed parameters.

Geom	Native (fun)	Alternate inputs
<code>geom_pdf()</code>	PDF	<code>cdf_fun</code> , <code>survival_fun</code> , <code>qf_fun</code> , <code>hf_fun</code>
<code>geom_cdf()</code>	CDF	<code>pdf_fun</code> , <code>survival_fun</code> , <code>qf_fun</code> , <code>hf_fun</code>
<code>geom_survival()</code>	Survival	<code>cdf_fun</code> , <code>pdf_fun</code> , <code>qf_fun</code>
<code>geom_qf()</code>	Quantile	<code>cdf_fun</code> , <code>pdf_fun</code> , <code>survival_fun</code>
<code>geom_hf()</code>	Hazard	<code>pdf_fun</code> , <code>cdf_fun</code> , <code>survival_fun</code> , <code>qf_fun</code>
<code>geom_chf()</code>	Cumulative hazard	<code>hf_fun</code> , <code>cdf_fun</code> , <code>survival_fun</code> , <code>pdf_fun</code> , <code>qf_fun</code>
<code>geom_cdf_discrete()</code>	CDF	<code>pmf_fun</code> , <code>survival_fun</code>
<code>geom_survival_discrete()</code>	Survival	<code>cdf_fun</code> , <code>pmf_fun</code>
<code>geom_qf_discrete()</code>	Quantile	<code>pmf_fun</code> , <code>cdf_fun</code> , <code>survival_fun</code>

Users whose primary interest is vector fields will find that **ggvfields** provides a richer collection of tools, including gradient fields, stream plots, and potential functions.

4 The probability distribution family

Statistics education depends heavily on the visual language of probability distributions. Instructors shade tail areas to illustrate p -values, draw CDFs to connect probabilities with quantiles, and mark central regions to delineate confidence intervals. **ggfunction** provides a family of geoms—univariate and bivariate, continuous and discrete—that map directly onto the standard functions associated with a probability distribution, each with built-in shading support. Throughout, the user supplies an R function (e.g., `dnorm`, `pnorm`, `qnorm`) and a range; **ggfunction** handles evaluation, shading, and validation automatically.

Each geom accepts its native function type through the `fun` parameter (e.g., a PDF for `geom_pdf()`, a CDF for `geom_cdf()`), but most also accept alternate function types that are converted internally. In each case the missing characterization is derived numerically: PDFs are obtained by differentiation or hazard integration, CDFs by integration or interpolation, survival functions by the identity $S = 1 - F$, and quantile functions by root-finding. Finite-support hazards can specify their integration origin with `hf_lower`. This cross-conversion feature lets users work with whichever characterization of a distribution is most natural—the PDF, CDF, survival function, quantile function, or hazard function—regardless of which geom they wish to plot. Exactly one function source must be provided; supplying more than one is an error. The bivariate layers `geom_pdf_2d()` and `geom_pmf_2d()`, introduced below, accept only their native function type via `fun`. Table 1 summarizes the available routes; the cross-conversion vignette documents the numerical methods and expected accuracy in detail.

4.1 Density functions: `geom_pdf()`

The probability density function (PDF) of a continuous random variable X satisfies




Figure 9: Three shading modes for `geom_pdf()`: lower tail ($p = 0.975$, left), central 95% interval (center), and two-tailed rejection region at $\alpha = 0.05$ (right).

$$f(x) \geq 0 \quad \text{and} \quad \int_{-\infty}^{\infty} f(x) dx = 1. \quad (2)$$

`geom_pdf()` evaluates a user-supplied density and renders it as a filled area with an overlaid outline. Alternatively, a CDF can be supplied via `cdf_fun` (differentiated numerically), a survival function via `survival_fun` (converted to a CDF via $F = 1 - S$ and then differentiated), a quantile function via `qf_fun` (inverted via interpolation to obtain a CDF, then differentiated), or a hazard function via `hf_fun` (converted directly using the identity $f(x) = h(x) \cdot e^{-H(x)}$). For finite-support hazards, `hf_lower` sets the lower integration limit for H . The underlying `StatPDF` validates the normalization property (2) using `stats::integrate()`, issuing a diagnostic via `cli` (Csárdi, 2024) if the integral departs from unity by more than a specified tolerance—a guard against accidentally passing an unnormalized function.

Four shading modes are supported, corresponding to common pedagogical operations.

Single threshold. The `p` parameter specifies a cumulative probability. When `lower.tail = TRUE` (the default), the region from the left boundary to the p -quantile is shaded, representing $P(X \leq x_p) = p$. Setting `lower.tail = FALSE` shades the complementary upper tail.

Two-sided interval. The `p_lower` and `p_upper` parameters define a central region—the natural representation for a $(1 - \alpha)$ confidence interval.

Tail shading. Setting `shade_outside = TRUE` inverts the two-sided region, shading both tails. This directly represents the rejection region of a two-sided hypothesis test at level α .

```
p1 <- ggplot() +
  geom_pdf(fun = dnorm, xlim = c(-3, 3), p = 0.975) +
  ggtitle("p = 0.975")

p2 <- ggplot() +
  geom_pdf(
    fun = dnorm, xlim = c(-3, 3),
    p_lower = 0.025, p_upper = 0.975
  ) + ggtitle("Central 95%")

p3 <- ggplot() +
  geom_pdf(
    fun = dnorm, xlim = c(-3, 3),
    p_lower = 0.025, p_upper = 0.975, shade_outside = TRUE
  ) + ggtitle(expression(alpha == 0.05 ~ "rejection region"))
```

p1 | p2 | p3

Highest density region. The `shade_hdr` parameter specifies a coverage probability and shades the corresponding highest density region (HDR)—the smallest subset of the domain containing the specified probability mass. For multimodal densities the HDR may be disconnected, producing multiple disjoint shaded intervals. The algorithm follows Otto and Kahle (2023): evaluate the density on the grid, normalize the values to sum to one, sort in descending order, accumulate until the target coverage is reached, and shade all grid intervals at or above the resulting density threshold.

```
f_mix <- function(x) 0.6 * dnorm(x, mean = -2, sd = 0.6) + 0.4 * dnorm(x, mean = 2, sd = 1.2)
```





Figure 10: The 80% highest density region of an asymmetric bimodal density ($0.6\mathcal{N}(-2, 0.6^2) + 0.4\mathcal{N}(2, 1.2^2)$), shading both modes as two disjoint intervals with more area allocated to the taller, narrower component.

```
ggplot() +
  geom_pdf(fun = f_mix, xlim = c(-5, 6), shade_hdr = 0.8)
```

The shading boundaries for the interval-based modes are determined by trapezoidal accumulation. For n evaluation points $x_1, \dots, x_n$ with density values $y_1, \dots, y_n$, the cumulative area at the k -th point is

$$A_k = \sum_{i=1}^{k-1} \frac{y_i + y_{i+1}}{2} (x_{i+1} - x_i). \quad (3)$$

The normalized values A_k/A_n are then compared against the specified probability thresholds to determine the shading boundaries.

4.2 Bivariate density functions: `geom_pdf_2d()`

The density family extends to bivariate distributions through `geom_pdf_2d()`, which visualizes a joint density $f(x, y)$ through its highest density regions (HDRs)—the smallest regions of the plane containing specified probability masses—or as a raw density raster. Rather than reimplement the HDR computation, the layer is a thin, probability-facing wrapper around `ggdensity` (Otto and Kahle, 2023): with type = "hdr" (the default), filled regions are drawn by `ggdensity::geom_hdr_fun()`, and with type = "hdr_lines" the region boundaries are drawn by `ggdensity::geom_hdr_lines_fun()`, so HDR computation, probability labeling, and default aesthetics remain with `ggdensity`. The density is supplied in `ggfunction`'s bivariate convention—`fun` receives a single vector $v = (x, y)$ and returns one density value, with parameters passed via `args`—and is adapted internally to `ggdensity`'s vectorized two-argument interface. The `probs` parameter sets the HDR coverage levels, defaulting to 99%, 95%, 80%, and 50%, and is ignored by type = "raster", which evaluates the density on a grid and renders a dark-gray raster with alpha mapped to the density value. For iso-density contours not calibrated to probability content, `geom_function_2d_1d()` remains the appropriate tool.

```
dbvn <- function(v, mu = c(0, 0), Sigma = diag(2)) {
  x <- matrix(v - mu, ncol = 1)
  Sinv <- solve(Sigma)
  1 / (2 * pi * sqrt(det(Sigma))) * exp(-0.5 * as.numeric(t(x) %*% Sinv %*% x))
}

p1 <- ggplot() +
  geom_pdf_2d(fun = dbvn, xlim = c(-3, 3), ylim = c(-3, 3),
    probs = c(0.5, 0.8, 0.95)) +
  coord_equal() + ggtitle('type = "hdr"')

Sigma <- matrix(c(1, 0.6, 0.6, 1), 2, 2)
p2 <- ggplot() +
```




Figure 11: Bivariate normal densities drawn by `geom_pdf_2d()`: filled 50/80/95% HDRs of the standard bivariate normal (left), HDR boundary lines for a correlated bivariate normal with covariance supplied via `args` (middle), and a raw density raster with alpha mapped to density (right).




Figure 12: The PMF of a Binomial(10,0.3) distribution, rendered as a lollipop chart.

```
geom_pdf_2d(fun = dbvn, args = list(Sigma = Sigma),
  xlim = c(-3, 3), ylim = c(-3, 3),
  probs = c(0.5, 0.8, 0.95), type = "hdr_lines") +
coord_equal() + ggtitle('type = "hdr_lines"')

p3 <- ggplot() +
  geom_pdf_2d(fun = dbvn, xlim = c(-3, 3), ylim = c(-3, 3),
    type = "raster") +
  coord_equal() + ggtitle('type = "raster"')

p1 | p2 | p3
```

4.3 Probability mass functions: `geom_pmf()`

For a discrete random variable with integer support, the PMF $p(k) = P(X = k)$ satisfies $\sum_k p(k) = 1$. `geom_pmf()` evaluates the PMF at each integer in `xlim` and renders the result as a lollipop chart—vertical segments from the x -axis to the probability value, capped with points. The underlying `StatPMF` validates normalization by summing the computed probabilities and issues a warning if they depart significantly from unity. When the support is not a sequence of consecutive integers, the support argument accepts an explicit numeric vector that overrides `xlim`.

```
ggplot() +
  geom_pmf(fun = dbinom, xlim = c(0, 10), args = list(size = 10, prob = 0.3))
```

The support argument enables distributions whose mass points are not consecutive integers. Figure~13 illustrates this with the exact distribution of the sample mean $\bar{X}_n = n^{-1} \sum_{i=1}^n X_i$ of $n = 10$ i.i.d. Bernoulli(0.3) draws. The support is $\{0, 0.1, 0.2, \dots, 1\}$ and $P(\bar{X}_n = k/n) = \binom{n}{k} 0.3^k 0.7^{n-k}$. The Central Limit Theorem approximation $\bar{X}_n \approx \mathcal{N}(0.3, 0.3 \times 0.7/10)$ is overlaid as a scaled density




Figure 13: Exact distribution of the sample mean of 10 i.i.d. Bernoulli(0.3) draws (lollipops) with the CLT normal approximation overlaid as a scaled density curve.

curve whose peak matches the tallest lollipop, facilitating direct visual comparison of shape without a secondary axis.

```
p <- 0.3; n <- 10
sd_mean <- sqrt(p * (1 - p) / n)
f_mean <- function(x) dbinom(round(x * n), size = n, prob = p)
max_pmf <- max(f_mean(seq(0, 1, by = 1/n)))
scale <- max_pmf / dnorm(p, mean = p, sd = sd_mean)
f_clt <- function(x) scale * dnorm(x, mean = p, sd = sd_mean)

ggplot() +
  geom_pmf(fun = f_mean, support = seq(0, 1, by = 1/n)) +
  geom_pdf(fun = f_clt, xlim = c(0, 1))
```

`geom_pmf()` supports the same shading modes as `geom_pdf()`. The `p` parameter shades lollipops up to the p -quantile; `p_lower` and `p_upper` define a two-sided interval; `shade_outside = TRUE` inverts it; and `shade_hdr` accepts one or more coverage levels, shading the smallest sets of mass points whose total probability meets or exceeds each target. Lollipops outside a p -based region are rendered at reduced opacity to preserve the distributional context, while `shade_hdr` maps each support point's smallest containing HDR to α as an ordered factor (`after_stat(probs)`), producing the same legend convention as the bivariate layers. Because a discrete distribution may not achieve the exact target coverages, `shade_hdr` finds the smallest HDR with coverage $\geq$ each requested level—a probability-threshold set $\{x : p(x) \geq c\}$ that includes all mass points tied at the cutoff, so ties are never split and the achieved coverage almost always overshoots the target—and issues a diagnostic via [cli](#) when an actual coverage differs from its target by more than 0.5 percentage points.

```
p1 <- ggplot() +
  geom_pmf(fun = dbinom, xlim = c(0, 10),
    args = list(size = 10, prob = 0.5), p = 0.8) +
  ggtitle("p = 0.8")

p2 <- ggplot() +
  geom_pmf(fun = dbinom, xlim = c(0, 10),
    args = list(size = 10, prob = 0.3, shade_hdr = c(0.5, 0.8, 0.95))) +
  ggtitle("shade_hdr = c(0.5, 0.8, 0.95)")

p1 | p2
```

4.4 Bivariate mass functions: `geom_pmf_2d()`

`geom_pmf_2d()` is the discrete counterpart: it evaluates a bivariate PMF over a lattice and renders the masses either as a balloon plot of points (`type = "point"`, the default, with size mapped to




Figure 14: Shading modes for `geom_pmf()`: the lower 80% by cumulative probability, with unshaded lollipops rendered at reduced opacity (left), and the 50/80/95% HDRs of a Binomial(10, 0.3) distribution, with each support point's smallest containing HDR mapped to alpha (right).

`after_stat(prob)`; `ggplot2::scale_size_area()` is recommended so that mass is proportional to point area) or as a heatmap of tiles (`type = "tile"`, with fill mapped to `after_stat(prob)`). The lattice is the product of per-axis supports—integers spanning $xlim \times ylim$, or exact values supplied via `support_x` and `support_y`—and the mass function follows the same `fun(c(x, y))` convention as `geom_pdf_2d()`. As in `geom_pmf()`, the total lattice mass is validated. The `shade_hdr` parameter accepts one or more coverage levels and highlights the corresponding highest density regions: each lattice point is assigned the smallest requested HDR that contains it, and the assignment is exposed as an ordered factor (`after_stat(probs)`) mapped to alpha by default, reproducing the legend convention of `geom_pdf_2d()`; points outside all requested regions render nearly transparent. Because a discrete distribution generally cannot achieve a target coverage exactly, the smallest HDR with coverage at least each target is used—with lattice points of exactly equal mass always assigned together, never split—and the actual, typically overshooting, coverages are reported. Distributions with non-product support are accommodated by evaluating over a bounding lattice and returning zero off the support: zero-mass cells are dropped by default (`drop_zeros = TRUE`), so a trinomial distribution, for example, renders only its simplex support.

```
dbinom2 <- function(v, sizes = c(10, 10), probs = c(0.5, 0.5)) {
  dbinom(v[1], sizes[1], probs[1]) * dbinom(v[2], sizes[2], probs[2])
}

p1 <- ggplot() +
  geom_pmf_2d(fun = dbinom2, xlim = c(0, 10), ylim = c(0, 10),
    args = list(probs = c(0.3, 0.7)), shade_hdr = c(0.5, 0.8, 0.95)) +
  scale_size_area() +
  ggtitle("Product binomial, 50/80/95% HDRs")

dtrinom <- function(v, size = 8, prob = c(0.3, 0.3, 0.4)) {
  if (sum(v) > size) return(0)
  dmultinom(c(v, size - sum(v)), prob = prob)
}

p2 <- ggplot() +
  geom_pmf_2d(fun = dtrinom, xlim = c(0, 8), ylim = c(0, 8),
    type = "tile") +
  coord_equal() + ggtitle("Trinomial on a simplex")

p1 | p2
```

4.5 Cumulative distribution and survival functions

The CDF $F(x) = P(X \leq x)$ and survival function $S(x) = 1 - F(x) = P(X > x)$ are complementary summaries of a distribution. **ggfunction** provides continuous and discrete variants of both.




Figure 15: Bivariate PMFs rendered by `geom_pmf_2d()`: the 50/80/95% highest density regions of a product $\text{Binomial}(10, 0.3) \times \text{Binomial}(10, 0.7)$ distribution in the default balloon rendering, with each lattice point's smallest containing HDR mapped to alpha (left), and a trinomial distribution with $n = 8$ evaluated over a bounding lattice in the tile rendering, showing only its simplex support (right).




Figure 16: The standard normal CDF.

Continuous CDF (`geom_cdf()`). `geom_cdf()` evaluates a user-supplied CDF (e.g., `pnorm`) and renders it as a line. Alternatively, a PDF can be supplied via `pdf_fun` (integrated numerically), a survival function via `survival_fun` (converted via $F = 1 - S$), a quantile function via `qf_fun` (inverted via interpolation), or a hazard function via `hf_fun` (the cumulative hazard is integrated from `hf_lower` and then $F = 1 - e^{-H}$). It supports the same `p`, `p_lower/p_upper`, and `shade_outside` shading parameters as `geom_pdf()`.

```
ggplot() +
  geom_cdf(fun = pnorm, xlim = c(-3, 3))
```

Discrete CDF (`geom_cdf_discrete()`). For discrete distributions, the CDF is a right-continuous step function. `geom_cdf_discrete()` accepts a CDF directly via `fun` (e.g., `pbinom`), a PMF via `pmf_fun` (from which the CDF is computed by cumulative summation), or a discrete survival function via `survival_fun` (converted via $F = 1 - S$), and renders the result with horizontal segments, dashed vertical jumps, open circles at the pre-jump value (the left limit), and closed circles at the achieved value. The `show_points` and `show_vert` parameters independently suppress the endpoint circles or vertical jump segments. The `support` argument accepts the full vector of mass points for distributions with non-integer or non-consecutive support; `xlim` filters the displayed range after cumulative quantities are computed.

```
ggplot() +
  geom_cdf_discrete(
    pmf_fun = dbinom, xlim = c(0, 10), args = list(size = 10, prob = 0.5)
  )
```

Continuous survival (`geom_survival()`). In reliability theory and biostatistics, $S(x) = 1 - F(x) = P(X > x)$ gives the probability that the event of interest has not yet occurred by time x .




Figure 17: The discrete CDF of a Binomial(10,0.5) distribution.




Figure 18: The survival function of an Exponential(0.5) distribution, $S(x) = e^{-0.5x}$.

`geom_survival()` accepts a survival function directly via `fun`, or derives S from a CDF via `cdf_fun` (computing $1 - F(x)$), a PDF via `pdf_fun` (integrating to obtain F , then computing $1 - F(x)$), or a quantile function via `qf_fun` (interpolating to obtain F , then computing $1 - F(x)$).

```
ggplot() +
  geom_survival(cdf_fun = pexp, xlim = c(0, 10), args = list(rate = 0.5))
```

Discrete survival (`geom_survival_discrete()`). For discrete distributions, $S(x) = 1 - F(x)$ is itself a right-continuous step function. `geom_survival_discrete()` accepts a discrete survival function directly via `fun`, or derives S from a CDF via `cdf_fun` or a PMF via `pmf_fun` (computing cumulative sums and then $1 - F$). It renders the result with the same visual conventions as `geom_cdf_discrete()`. PMF-derived survival curves use `support` as the full computational support and `xlim` as a display filter.

```
ggplot() +
  geom_survival_discrete(
    pmf_fun = dbinom, xlim = c(0, 10), args = list(size = 10, prob = 0.5)
  )
```

4.6 Quantile functions

The quantile function $Q(p) = \inf\{x : F(x) \geq p\}$ inverts the CDF. **ggfunction** provides continuous and discrete variants.

Continuous (`geom_qf()`). `geom_qf()` evaluates a user-supplied quantile function (e.g., `qnorm`) over the unit interval $(0, 1)$. Alternatively, a CDF can be supplied via `cdf_fun` (inverted numerically by




Figure 19: The discrete survival function of a Binomial(10, 0.5) distribution, descending from 1 to 0.




Figure 20: The quantile function of the standard normal distribution.

root-finding), a PDF via `pdf_fun` (integrated to obtain the CDF, then inverted), or a survival function via `survival_fun` (converted to a CDF via $F = 1 - S$, then inverted). Evaluation points are placed at Chebyshev nodes of the first kind, $p_k = (1 - \cos((2k - 1)\pi/2n))/2$ for $k = 1, \dots, n$, which cluster near 0 and 1 where quantile functions are typically most curved and avoid evaluating at the exact endpoints, preventing $\pm\infty$ for unbounded distributions.

```
ggplot() +
  geom_qf(fun = qnorm)
```

Discrete (`geom_qf_discrete()`). The quantile function of a discrete distribution is a left-continuous step function on $[0, 1]$. `geom_qf_discrete()` accepts a quantile function directly via `fun` (e.g., `qbinom`), or derives it from a PMF via `pmf_fun` (computing cumulative sums over the full support and inverting), a CDF via `cdf_fun` (evaluated on the integer support and inverted), or a discrete survival function via `survival_fun` (converted to a CDF via $F = 1 - S$ and then inverted). It renders the result with horizontal segments on $[0, 1]$, dashed vertical jumps, closed circles at the bottom of each jump (the value is achieved), and open circles at the top (the next value is not yet reached). The support argument behaves as in `geom_cdf_discrete()`.

```
ggplot() +
  geom_qf_discrete(
    pmf_fun = dbinom, xlim = c(0, 10), args = list(size = 10, prob = 0.5)
  )
```

4.7 Hazard functions: `geom_hf()`

The hazard function




Figure 21: The discrete quantile function of a Binomial(10, 0.5) distribution as a left-continuous step function on $[0, 1]$.

$$h(x) = \frac{f(x)}{S(x)} = \frac{f(x)}{1 - F(x)} \quad (4)$$

represents the instantaneous rate of failure at time x , conditional on survival to that point (Casella and Berger, 2002). `geom_hf()` accepts a hazard function directly via `fun`, or derives h from a PDF and CDF supplied via `pdf_fun` and `cdf_fun`. When only one of `pdf_fun` or `cdf_fun` is provided, the missing component is derived numerically (by integration or differentiation, respectively). A survival function can also be supplied via `survival_fun` (converted to a CDF via $F = 1 - S$, then differentiated to obtain f), or a quantile function via `qf_fun` (interpolated to obtain F , then differentiated). The `args` parameter applies to all supplied functions; `pdf_args` and `cdf_args` provide overrides where the PDF and CDF require different parameterizations.

Division by $S(x)$ can cause numerical instability in the tail as $S(x) \rightarrow 0$. `StatHF` guards against this by replacing values where $S(x) \leq 0$ with `NaN`, and `GeomHF` filters these before rendering.

Three canonical hazard shapes arise in reliability modeling: decreasing (infant mortality, where early failures dominate), constant (the memoryless Exponential distribution), and increasing (wear-out failure, where older units fail at higher rates).

```
p_decr <- ggplot() +
  geom_hf(
    pdf_fun = dweibull, cdf_fun = pweibull,
    xlim = c(0.01, 5), args = list(shape = 0.5, scale = 2)
  ) + ggtitle("Decreasing (Weibull)")
```

```
p_flat <- ggplot() +
  geom_hf(
    pdf_fun = dexp, cdf_fun = pexp,
    xlim = c(0.01, 10), args = list(rate = 0.5)
  ) + ggtitle("Flat (Exponential)")
```

```
p_incr <- ggplot() +
  geom_hf(
    pdf_fun = dnorm, cdf_fun = pnorm,
    xlim = c(-3, 3), args = list(mean = 0, sd = 1)
  ) + ggtitle("Increasing (Normal)")
```

```
p_decr | p_flat | p_incr
```

4.8 Cumulative hazard functions: `geom_chf()`

The cumulative hazard function

$$H(x) = \int_{\ell}^x h(t) dt = -\log S(x) \quad (5)$$




Figure 22: Three canonical hazard shapes: decreasing (Weibull(shape=0.5, scale=2), left), constant (Exponential(0.5), center), and increasing ($\mathcal{N}(0, 1)$, right).




Figure 23: Cumulative hazard functions for three distributions: Exponential(0.5) (left, linear), Weibull(2, 1) (center, convex), and $\mathcal{N}(0, 1)$ (right, eventually superlinear). For the exponential, $H(x) = 0.5x$ confirms the constant hazard rate.

accumulates the instantaneous hazard over time. It is a monotonically non-decreasing function starting at zero that rises without bound for distributions with unbounded support. In reliability analysis, $H(x)$ is often more convenient than $h(x)$ because it is monotone and avoids the division by $S(x)$ that can cause numerical instability in the hazard function's tail.

`geom_chf()` renders $H(x)$ as a curve. It accepts a cumulative hazard function directly via `fun`, or derives H from any other standard characterization: a hazard function via `hf_fun` (integrated numerically from `hf_lower`), a CDF via `cdf_fun` ($H = -\log(1 - F)$), a survival function via `survival_fun` ($H = -\log S$), a PDF via `pdf_fun` (integrated to obtain F , then transformed), or a quantile function via `qf_fun` (interpolated to obtain F , then transformed).

```
p1 <- ggplot() +
  geom_chf(cdf_fun = pexp, xlim = c(0, 10), args = list(rate = 0.5)) +
  ggtitle("Exponential(0.5)")

p2 <- ggplot() +
  geom_chf(cdf_fun = pweibull, xlim = c(0, 3), args = list(shape = 2, scale = 1)) +
  ggtitle("Weibull(2, 1)")

p3 <- ggplot() +
  geom_chf(cdf_fun = pnorm, xlim = c(-3, 3)) +
  ggtitle("Normal(0, 1)")

p1 | p2 | p3
```

5 The empirical data family

Data analysis frequently requires visualizing not a theoretical distribution but the distribution of an observed sample. **ggfunction**'s empirical data family applies the same visual language as the




Figure 24: Empirical CDFs for two groups from populations $\mathcal{N}(0, 1)$ and $\mathcal{N}(2, 1)$, with simultaneous 95% DKW confidence bands. The bands are narrower for larger samples.

probability distribution family to sample data. The empirical CDF, quantile function, PMF, and cumulative hazard geoms tabulate the empirical distribution via a shared helper that places mass c_k/n at each distinct observed value x_k , where c_k is the count of occurrences and n is the total sample size; ties are handled correctly. The PP, QQ, and SP diagnostic geoms use the related order-statistic representation of the same sample. Because the geoms are standard `ggplot2` layers, multiple groups can be overlaid by mapping the colour aesthetic in the enclosing `ggplot()` call.

5.1 Empirical CDF: `geom_ecdf()`

The empirical cumulative distribution function

$$\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(X_i \leq x) \quad (6)$$

is the natural nonparametric estimate of the true CDF. `geom_ecdf()` renders $\hat{F}_n$ as a right-continuous step function—horizontal segments, dashed vertical jumps, and open/closed endpoint circles—matching the visual conventions of `geom_cdf_discrete()`. When the sample is large (more than 50 distinct values), the vertical jump segments and endpoint circles are suppressed automatically to avoid visual clutter; this threshold can be overridden by setting `show_vert` and `show_points` explicitly.

An optional simultaneous DKW/Massart confidence band is drawn around the step function by default. The DKW inequality (Massart, 1990) is a finite-sample result that implies

$$P\left(\sup_x |\hat{F}_n(x) - F(x)| \leq \varepsilon\right) \geq 1 - \alpha, \quad \varepsilon = \sqrt{\frac{\log(2/\alpha)}{2n}}, \quad (7)$$

so the ribbon $[\hat{F}_n(x) - \varepsilon, \hat{F}_n(x) + \varepsilon]$ (clipped to $[0, 1]$) covers the true CDF F simultaneously at all x with probability at least $1 - \alpha$. The confidence level is controlled by the `level` parameter (default 0.95), and the ribbon's transparency by `conf_alpha` (default 0.3).

```
set.seed(1)
df_ecdf <- data.frame(
  x      = c(rnorm(60), rnorm(40, mean = 2)),
  group  = rep(c("N(0,1), n=60", "N(2,1), n=40"), c(60, 40))
)

ggplot(df_ecdf, aes(x = x, colour = group)) +
  geom_ecdf(show_points = FALSE, show_vert = FALSE) +
  labs(x = "x", y = expression(hat(F)[n](x)), colour = NULL)
```

5.2 Empirical quantile function: `geom_eqf()`

The empirical quantile function is the left-continuous inverse of the empirical CDF:

$$Q_n(p) = \inf\{x : \hat{F}_n(x) \geq p\}, \quad p \in (0, 1). \quad (8)$$

`geom_eqf()` renders Q_n as a left-continuous step function on $[0, 1]$ using the same visual conventions as `geom_qf_discrete()`. The same auto-suppression rule applies: jump segments and endpoint circles are hidden when the sample has more than 50 distinct values.

The confidence band for the quantile function follows directly from inverting (7): since $\hat{F}_n(x) - \varepsilon \leq F(x) \leq \hat{F}_n(x) + \varepsilon$ simultaneously, applying Q to all three sides yields $Q_n(p - \varepsilon) \leq Q(p) \leq Q_n(p + \varepsilon)$.

This same probability-scale band is the basis for the PP, QQ, and SP diagnostic layers.

5.3 PP, QQ, and SP diagnostic plots: `geom_ppplot()`, `geom_qqplot()`, and `geom_spplot()`

Suppose $X_1, \dots, X_n$ are to be compared with a fully specified null distribution F_0 , with quantile function $Q_0 = F_0^{\leftarrow}$. Let $x_{(1)} \leq \dots \leq x_{(n)}$ denote the order statistics and define plotting positions

$$p_i = \frac{i - 1/2}{n}, \quad i = 1, \dots, n. \quad (9)$$

A probability-probability (PP) plot displays the null probabilities $F_0(x_{(i)})$ against p_i . A quantile-quantile (QQ) plot displays the observed quantiles $x_{(i)}$ against the null quantiles $Q_0(p_i)$. A stabilized probability (SP) plot applies the arcsine-square-root variance-stabilizing transform

$$g(p) = \frac{2}{\pi} \sin^{-1} \sqrt{p} \quad (10)$$

to both PP coordinates, displaying $g\{F_0(x_{(i)})\}$ against $g(p_i)$; under the null the transformed points have approximately equal variance across the entire unit interval, so departures in the tails are as visually prominent as departures in the middle. In each case, agreement with the null distribution corresponds to points falling close to the identity line. Michael (1983) discusses acceptance regions for PP, QQ, and stabilized probability plots. The PP and SP layers expect a CDF by default, while the QQ layer expects a quantile function; alternate function sources are converted internally using the same CDF/QF helpers used by the theoretical distribution geoms.

Under $H_0 : F = F_0$, the probability integral transform gives $F_0(X_i) \stackrel{\text{iid}}{\sim} \text{Unif}(0, 1)$. Therefore (7) gives a simultaneous DKW band on the probability scale. For PP plots, the ribbon is simply

$$\max(0, p_i - \varepsilon) \leq F_0(x_{(i)}) \leq \min(1, p_i + \varepsilon). \quad (11)$$

For SP plots, the monotone transform g of (10) is applied to both sides of (11), carrying the same simultaneous band to the stabilized scale. For QQ plots, the probability limits are transformed back to the data scale:

$$Q_0\{\max(0, p_i - \varepsilon)\} \leq x_{(i)} \leq Q_0\{\min(1, p_i + \varepsilon)\}. \quad (12)$$

Thus the three diagnostics show the same comparison on different scales: PP bands are parallel probability bands, QQ bands widen or narrow according to the shape of Q_0 , and SP bands carry the parallel probability bands to the variance-stabilized scale. The bands are finite-sample and distribution-free for a fully specified null. When parameters are estimated from the data, the plots remain useful visual diagnostics, but the DKW bands are conservative rather than Lilliefors-adjusted.

```
set.seed(3)
```

```
df_normal <- data.frame("x" = rnorm(50))
df_exp <- data.frame("x" = rexp(50) - 1)
```

```
ppqq_theme <- theme_gray(base_size = 8) +
  theme(
    plot.title = element_text(size = 8),
    axis.title = element_text(size = 8),
    axis.text = element_text(size = 7)
  )
```

```
p_norm_pp <- ggplot(df_normal, aes(x = x)) +
  geom_ppplot(fun = pnorm, colour = "black", size = 1.2) +
  coord_equal() +
```




Figure 25: PP, QQ, and SP diagnostics for samples compared with a fully specified $\mathcal{N}(0, 1)$ null. The normal sample (top row) follows the identity line and remains within the DKW bands, whereas the shifted exponential sample (bottom row) shows the curvature and tail departures expected from skewness; the SP panels show the same comparison on the variance-stabilized scale.

```
ggtitle("Normal data: PP")

p_norm_qq <- ggplot(df_normal, aes(x = x)) +
  geom_qqplot(fun = qnorm, colour = "black", size = 1.2) +
  coord_equal() +
  ggtitle("Normal data: QQ")

p_norm_sp <- ggplot(df_normal, aes(x = x)) +
  geom_spplot(fun = pnorm, colour = "black", size = 1.2) +
  coord_equal() +
  ggtitle("Normal data: SP")

p_exp_pp <- ggplot(df_exp, aes(x = x)) +
  geom_ppplot(fun = pnorm, colour = "black", size = 1.2) +
  coord_equal() +
  ggtitle("Exponential data: PP")

p_exp_qq <- ggplot(df_exp, aes(x = x)) +
  geom_qqplot(fun = qnorm, colour = "black", size = 1.2) +
  coord_equal() +
  ggtitle("Exponential data: QQ")

p_exp_sp <- ggplot(df_exp, aes(x = x)) +
  geom_spplot(fun = pnorm, colour = "black", size = 1.2) +
  coord_equal() +
  ggtitle("Exponential data: SP")

(p_norm_pp | p_norm_qq | p_norm_sp) / (p_exp_pp | p_exp_qq | p_exp_sp) & ppqq_theme
```

5.4 Empirical PMF: `geom_epmf()`

For observed data that are discrete or treated as such, the empirical PMF places mass c_k/n at each distinct observed value x_k and renders the result as a lollipop chart using the same conventions as `geom_pmf()`.

```
set.seed(2)
df_pmf <- data.frame(
  x      = round(c(rnorm(40), rnorm(40, mean = 2)), 1),
```




Figure 26: Empirical PMFs for two groups of 40 observations each from $\mathcal{N}(0, 1)$ and $\mathcal{N}(2, 1)$, colored by group.

```
group = rep(c("A", "B"), each = 40)
)

ggplot(df_pmf, aes(x = x, colour = group)) +
  geom_pmf() +
  labs(x = "x", y = "Empirical probability", colour = "Group")
```

5.5 Empirical cumulative hazard: `geom_echf()`

The empirical cumulative hazard function transforms the ECDF to the hazard scale:

$$\hat{H}_n(x) = -\log(1 - \hat{F}_n(x)). \quad (13)$$

At the largest observation, $\hat{F}_n(x) = 1$ and $\hat{H}_n(x) = \infty$; this boundary is dropped automatically and the step function extends to the right panel edge at the last finite value.

The simultaneous confidence band is obtained by applying the monotone transformation $H = -\log(1 - F)$ to the DKW bounds (7). Because the DKW inequality guarantees $\hat{F}_n(x) - \varepsilon \leq F(x) \leq \hat{F}_n(x) + \varepsilon$ simultaneously at all x , applying $H = -\log(1 - F)$ to both sides (and clipping the CDF bounds to $[0, 1)$) yields a simultaneous band for $H(x)$ with the same coverage guarantee.

```
set.seed(5)

df_correct <- data.frame(x = rexp(80, rate = 0.5))
p_correct <- ggplot(df_correct, aes(x = x)) +
  geom_echf(show_points = FALSE, show_vert = FALSE) +
  geom_chf(cdf_fun = pexp, xlim = c(0, max(df_correct$x)),
    args = list(rate = 0.5), colour = "red") +
  labs(x = "x", y = expression(hat(H)[n](x))) +
  ggtitle("Correct model") +
  theme_minimal()

df_wrong <- data.frame(x = rexp(80, rate = 1))
p_wrong <- ggplot(df_wrong, aes(x = x)) +
  geom_echf(show_points = FALSE, show_vert = FALSE) +
  geom_chf(cdf_fun = pexp, xlim = c(0, max(df_wrong$x)),
    args = list(rate = 0.5), colour = "red") +
  labs(x = "x", y = expression(hat(H)[n](x))) +
  ggtitle("Wrong model") +
  theme_minimal()

p_correct | p_wrong
```




Figure 27: Empirical cumulative hazard functions for simulated exponential data. An $\text{Exp}(0.5)$ sample (left) is overlaid with the theoretical $H(x) = 0.5x$ (red line), which lies within the 95% confidence band. An $\text{Exp}(1)$ sample (right) is tested against the wrong model $H(x) = 0.5x$, which departs from the band.

The theoretical cumulative hazard (red) threads through the band for the correct model (left panel) but departs visibly for the misspecified model (right panel), where the true rate is twice the hypothesized value.

5.6 Kaplan–Meier survival curves: `geom_ecdf_km()`

The empirical geoms above assume complete (uncensored) data. In practice, survival data are often right-censored: some subjects leave the study before the event is observed. The Kaplan–Meier product-limit estimator ([Kaplan and Meier, 1958](#)) handles this by adjusting the risk set at each event time:

$$\hat{S}(t) = \prod_{t_j \leq t} \left(1 - \frac{d_j}{n_j}\right), \quad (14)$$

where d_j is the number of events at time t_j and n_j is the number at risk just before t_j . `geom_ecdf_km()` renders this as a decreasing step function, with a simultaneous equal-precision Greenwood confidence band ([Greenwood, 1926](#); [Nair, 1984](#)) and censoring marks at observations where the event was not observed. The data require two aesthetics: `x` (the observed time) and `status` (1 = event, 0 = censored).

```
set.seed(42)
n_km <- 60
true_time <- rexp(n_km, rate = 0.5)
cens_time <- rexp(n_km, rate = 0.2)
df_km <- data.frame(
  time = pmin(true_time, cens_time),
  status = as.integer(true_time <= cens_time)
)

p_km <- ggplot(df_km, aes(x = time, status = status)) +
  geom_ecdf_km() +
  labs(x = "Time", y = expression(hat(S)(t))) +
  ggtitle("Kaplan-Meier estimate") +
  theme_minimal()

p_km_gof <- ggplot(df_km, aes(x = time, status = status)) +
  geom_ecdf_km(show_points = FALSE, show_vert = FALSE) +
  geom_survival(cdf_fun = pexp, xlim = c(0, max(df_km$time)),
    args = list(rate = 0.5), colour = "red") +
  labs(x = "Time", y = expression(hat(S)(t))) +
  ggtitle("KM vs. Exp(0.5) theory") +
  theme_minimal()
```




Figure 28: Kaplan–Meier survival curves for simulated censored data. Event times follow $\text{Exp}(0.5)$ and censoring times follow $\text{Exp}(0.2)$. The left panel shows the KM estimate with a Greenwood 95% simultaneous equal-precision band and censoring marks; the right panel overlays the true survival curve (red).

```
p_km | p_km_gof
```

The theoretical survival curve (red) lies within the Greenwood EP confidence band, confirming consistency with the $\text{Exp}(0.5)$ model.

5.7 Nelson–Aalen cumulative hazard: `geom_echf_na()`

The Nelson–Aalen estimator (Nelson, 1972; Aalen, 1978) is the censored-data analogue of the empirical cumulative hazard:

$$\hat{H}(t) = \sum_{t_j \leq t} \frac{d_j}{n_j}, \quad (15)$$

with variance estimated by $\widehat{\text{Var}}[\hat{H}(t)] = \sum_{t_j \leq t} d_j / n_j^2$. `geom_echf_na()` renders this as an increasing step function with a pointwise normal confidence band.

```
p_na <- ggplot(df_km, aes(x = time, status = status)) +
  geom_echf_na() +
  labs(x = "Time", y = expression(hat(H)(t))) +
  ggtitle("Nelson-Aalen estimate") +
  theme_minimal()

p_na_gof <- ggplot(df_km, aes(x = time, status = status)) +
  geom_echf_na(show_points = FALSE, show_vert = FALSE) +
  geom_chf(cdf_fun = pexp, xlim = c(0, max(df_km$time)),
    args = list(rate = 0.5), colour = "red") +
  labs(x = "Time", y = expression(hat(H)(t))) +
  ggtitle("NA vs. Exp(0.5) theory") +
  theme_minimal()
```

```
p_na | p_na_gof
```

The theoretical cumulative hazard (red) threads through the confidence band, consistent with the generating model.

5.8 Validity of the confidence band

The DKW inequality (7) is a *finite-sample* result: it holds for every $n \geq 1$, not merely asymptotically. It guarantees that the simultaneous confidence band contains the true CDF everywhere with probability




Figure 29: Nelson–Aalen cumulative hazard estimates. The same censored data as in the KM figure. The left panel shows the Nelson–Aalen estimate with a 95% pointwise normal confidence band; the right panel overlays the theoretical $H(x) = 0.5x$ (red).

at least $1 - \alpha$ at any sample size. Because the KS statistic $D_n = \sup_x |\hat{F}_n(x) - F(x)|$ has a distribution-free null distribution for continuous F , the coverage guarantee holds regardless of the shape of the true distribution.

We verify the guarantee by simulation. For each combination of sample size n and nominal level $1 - \alpha$, we draw $B = 2,000$ independent samples from $\mathcal{N}(0, 1)$, compute D_n via `ks.test()`, and record the fraction of replications for which $D_n \leq \varepsilon_n$, where ε_n is the half-width used by `geom_ecdf()`.

```
set.seed(20240101)
B <- 10000
ns <- c(10, 20, 50, 100, 200, 500, 1000)
levels <- c(0.90, 0.95, 0.99)
eps_fn <- function(n, lv) sqrt(log(2 / (1 - lv))) / (2 * n))

sim_res <- do.call(rbind, lapply(ns, function(n) {
  do.call(rbind, lapply(levels, function(lv) {
    eps <- eps_fn(n, lv)
    dn <- replicate(B,
      suppressWarnings(ks.test(rnorm(n), "pnorm", exact = FALSE)$statistic))
    data.frame(n = n, level = lv, empirical = mean(dn <= eps))
  }))
}))

ggplot(sim_res, aes(x = n, y = empirical,
  colour = factor(level), group = factor(level))) +
  geom_hline(aes(yintercept = level, colour = factor(level)),
    linetype = "dashed", linewidth = 0.4) +
  geom_line(linewidth = 0.8) +
  geom_point(size = 2) +
  scale_x_log10(breaks = ns) +
  scale_y_continuous(
    labels = scales::percent_format(accuracy = 0.1),
    limits = c(0.88, 1.0)
  ) +
  scale_colour_manual(
    values = c("0.9" = "steelblue", "0.95" = "firebrick", "0.99" = "forestgreen"),
    labels = c("0.9" = "90%", "0.95" = "95%", "0.99" = "99%"),
    name = "Nominal"
  ) +
  labs(x = "Sample size n (log scale)", y = "Empirical coverage") +
  theme_minimal()
```

Empirical coverage (solid lines) lies above each dashed nominal threshold across all tested sample sizes, confirming that the DKW bound is a *finite-sample* guarantee: it holds for every $n \geq 1$, not merely




Figure 30: Empirical simultaneous coverage of the KS confidence band over 10,000 simulations from $\mathcal{N}(0,1)$. Dashed lines show the nominal levels; solid curves show empirical coverage. Coverage is everywhere at or above nominal for all n , confirming the finite-sample validity of the DKW bound (Massart, 1990). The slight conservatism at small n diminishes as n grows because the bound’s constant is asymptotically tight.

in the limit. The slight conservatism at small n reflects that the bound is an inequality; the constant $C = 2$ is tight (Massart, 1990) in the sense that it cannot be universally reduced, and the conservatism diminishes as n grows because the KS distribution converges to the Kolmogorov distribution.

6 Implementation

6.1 Architecture

Each **ggfunction** layer is implemented as a pair of ggproto objects—a Stat and a Geom—wired together by a constructor function. The constructor creates a standard **ggplot2** `layer()` call, threading the user-supplied function and its parameters through to the stat’s `compute_group()` method. That method performs the core computation: evaluating the function on a grid or sequence and returning a data frame whose columns correspond to the aesthetic variables expected by the geom.

The separation of computation (stat) from rendering (geom) is not merely organizational. It allows users to substitute alternative geoms or compose multiple layers from the same stat. The `StatFunction2d` object, for instance, computes a grid of scalar-field values that can be rendered by `GeomRaster`, `GeomContour`, or `GeomContourFilled`, selected at construction time through the type parameter.

6.2 Function injection

R’s `...` mechanism does not compose well when a function is called inside a stat’s `compute_group()` method, which must simultaneously receive **ggplot2**-internal parameters. **ggfunction** adopts the approach used by `ggplot2::stat_function()`: additional arguments are collected in a named list (`args`) and injected at the call site via `rlang::inject()`:

```
fun_injected <- function(x) {
  rlang::inject(fun(x, !!!args))
}
```

This pattern is simple, composable, and avoids the ambiguity of routing function arguments through `...`.

6.3 Numerical considerations

Several layers involve numerical integration or accumulation:

- **PDF validation** uses `stats::integrate()` with adaptive quadrature to verify that the supplied function integrates to one over the specified domain.

- **Cross-conversion** between function types is handled by internal helpers. `pdf_to_cdf()` obtains a CDF by numerically integrating the PDF via `stats::integrate()` at each evaluation point. `cdf_to_pdf()` obtains a PDF by central finite differences of the CDF, $(F(x+h) - F(x-h))/2h$ with $h = 10^{-5}$. `cdf_to_qf()` inverts the CDF by root-finding via `stats::uniroot()`, with adaptive bound widening to accommodate distributions whose support is not known in advance. `hf_to_cdf()` integrates a hazard function from `hf_lower` to obtain the cumulative hazard $H(x)$ and returns $F(x) = 1 - e^{-H(x)}$. `hf_to_pdf()` uses the identity $f(x) = h(x) \cdot e^{-H(x)}$ directly, avoiding the numerical instability that would arise from chaining `hf_to_cdf` with `cdf_to_pdf`. These conversions are approximate; for distributions where analytic forms are available, supplying the native function type directly will be more accurate.
- **Shading boundaries** are computed by trapezoidal accumulation (3), which is efficient and sufficient for the smooth densities encountered in practice.
- **Boundary interpolation** in `GeomFunction1d` uses `stats::approxfun()` to compute exact y -values at shade endpoints that fall between evaluation points.
- **Quantile function evaluation** in `geom_qf()` uses Chebyshev nodes of the first kind rather than a uniform grid, concentrating evaluation points near $p = 0$ and $p = 1$ where the function is most curved and avoiding evaluation at the exact endpoints.
- **PP, QQ, and SP diagnostics** sort finite observations, compute plotting positions with `stats::ppoints(n, a = 1/2)`, and derive their DKW bands from the same half-width as `geom_ecdf()`. The QQ ribbon transforms the probability-scale band through the supplied or derived quantile function, and the SP layer applies the arcsine-square-root transform (10) to both the plotted coordinates and the band limits.
- **Vector field integration** (via `ggvfields`) uses a fourth-order Runge–Kutta method with configurable step size and maximum iterations.

The default resolution of $n = 101$ for one-dimensional functions and $n = 50$ for two-dimensional grids (yielding 2500 evaluation points) balances visual fidelity against computational cost. Users may increase n for functions with fine-scale structure.

7 Composability and the grammar and limitations

Because every **ggfunction** layer is a standard **ggplot2** layer, it composes freely with the rest of the grammar. Layers can be overlaid, scaled, themed, annotated, and faceted exactly as one would with any other geom. The following example superimposes two normal densities with different means and spreads, demonstrating that function visualization and standard **ggplot2** idioms are fully interoperable:

```
ggplot() +
  geom_pdf(
    fun = dnorm, xlim = c(-5, 8),
    args = list(mean = 0, sd = 1), alpha = 0.4
  ) +
  geom_pdf(
    fun = dnorm, xlim = c(-5, 8),
    args = list(mean = 3, sd = 1.5), alpha = 0.4
  ) +
  labs(x = "x", y = "f(x)", title = "Comparing two normal densities") +
  theme_minimal()
```

This composability is the primary advantage of embedding function visualization within the grammar of graphics, rather than providing standalone plotting functions.

While **ggfunction** composes naturally with the grammar, it is worth noting what it does *not* do: it does not extend the grammar of graphics itself. In Wilkinson’s formulation and its **ggplot2** implementation, aesthetics are visual properties of geometric objects—position, color, size, shape, and so on. A function such as $f(x)$ is not such a property; rather, it *references* aesthetics, mapping one visual dimension to another. It is therefore unclear whether functions can serve as aesthetics in any principled sense, since they occupy a different conceptual role than the data-to-visual mappings the grammar was designed around.

This distinction has practical consequences. R is a vectorized language, but not everything in R is a vector of elementary objects, and functions form the prime example. A function cannot be a column of a data frame, and it cannot be mapped within an `aes()` call. As a result, some grammar operations that users take for granted—faceting across levels of a variable, for instance—do not




Figure 31: Two normal densities with different means and spreads overlaid in a single plot, demonstrating composability with the `ggplot2` grammar of graphics.

transfer directly to function visualization. **ggfunction** works within these constraints by accepting function objects as layer-level parameters rather than as mapped aesthetics, an approach that is plainly useful but inherently limited. This suggests that further research into the grammar of graphics, and correspondingly into the implementation of **ggplot2** or an extension package, could enable richer treatment of functions as first-class objects in statistical visualization.

8 Related packages

Several existing packages overlap with **ggfunction** in purpose, and understanding their scope helps clarify where **ggfunction** fits.

1. **ggplot2** (Wickham, 2016). The `stat_function()` layer handles scalar functions $\mathbb{R} \rightarrow \mathbb{R}$ competently, and `geom_function_1d_1d()` builds directly on this foundation. **ggplot2** provides no analogues for parametric curves, scalar fields, vector fields, or probability-specific shading operations.
2. **ggdist** (Kay, 2024). This package excels at visualizing distributional uncertainty in Bayesian workflows: posterior distributions, predictive intervals, and uncertainty estimates from fitted models. Its geoms work primarily with samples or distributional objects rather than the $d/p/q/h$ function families that **ggfunction** targets.
3. **ggdensity** (Otto and Kahle, 2023). This package improves bivariate density visualization in **ggplot2**, with a particular focus on highest density regions (HDRs). **ggfunction**'s `shade_hdr` parameter in `geom_pdf()` follows the HDR computation approach introduced there: normalizing density values, sorting in descending order, and accumulating until the target coverage is reached. More directly, `geom_pdf_2d()` delegates HDR computation and rendering to **ggdensity**'s `geom_hdr_fun()` and `geom_hdr_lines_fun()`, adapting the `fun(c(x, y))` convention shared by **ggfunction**'s bivariate layers to **ggdensity**'s vectorized two-argument interface; its raster mode evaluates the same bivariate density convention directly on a grid.
4. **mosaic** (Pruim et al., 2017). The `plotDist()` function provides quick plots of named distributions and supports shading, but it is built on the **lattice** (Sarkar, 2008) graphics framework rather than **ggplot2**, and therefore cannot be composed with **ggplot2** layers.
5. **metR** (Campitelli, 2024). This package provides contour and vector-field geoms designed for meteorological data, requiring precomputed grids as inputs rather than function objects. Users who already have gridded data may find **metR** more convenient; users who want to visualize a function directly will find **ggfunction** more natural.
6. **ggvfields** (Turner et al., 2026). **ggfunction**'s `geom_function_2d_2d()` delegates to **ggvfields** for streamline computation. For users whose primary interest is vector fields—including stream plots, gradient fields, and potential functions—**ggvfields** offers a substantially richer collection of tools built on the same **ggplot2** extension architecture.

Taken together, these packages reflect the breadth of interest in function visualization within the R ecosystem. **ggfunction** is distinguished by its unified dimensional taxonomy and its integration of the full probability distribution function family—capabilities that, to our knowledge, no other single package provides within the grammar of graphics.

9 Summary

ggfunction extends **ggplot2** with a principled framework for visualizing mathematical functions, theoretical probability distributions, and empirical data distributions. Its organizing taxonomy—classifying functions by the dimensions of their domain and codomain—yields a small, memorable API that covers the function types most commonly encountered in mathematics and statistics. The probability distribution family spans continuous and discrete distributions—density, CDF, PMF, quantile, discrete CDF, discrete quantile, discrete survival, survival, hazard, and cumulative hazard functions—and extends to bivariate densities and mass functions through highest density regions and lattice displays. Each supports the shading operations central to statistical pedagogy—marking quantiles, delineating confidence regions, highlighting rejection areas, and shading highest density regions. The empirical data family complements these theoretical geoms by computing the empirical CDF, quantile function, PP, QQ, and SP diagnostic plots, PMF, and cumulative hazard function directly from samples, each paired with a simultaneous DKW/Massart confidence band where appropriate.

By implementing each layer as a ggproto stat-geom pair, the package inherits the full composability of the grammar of graphics. Functions and empirical summaries can be overlaid, themed, faceted, and annotated using the same tools that **ggplot2** users already know. The args injection pattern provides a clean interface for parameterized function families, and built-in validation catches common errors—such as unnormalized densities—before they produce misleading graphics.

ggfunction is available at <https://github.com/dusty-turner/ggfunction> and can be installed via `pak::pak("dusty-turner/ggfunction")`.

10 Acknowledgments

This package and manuscript were made possible by agentic coding platforms, primarily OpenAI's Codex (GPT 5.3 Codex Extra High) and Anthropic's Claude (Sonnet and Opus 4.6). Early drafts of stats and geoms associated with CDFs, in particular the empirical CDF, were done by James Otto and were his contribution to the work. The package and manuscript were written by Dusty Turner and David Kahle, and early reflections on the design of the package were provided by Rodney Sturdivant.

References

- O. Aalen. Nonparametric inference for a family of counting processes. *The Annals of Statistics*, 6(4): 701–726, 1978. doi: 10.1214/aos/1176344247. [p1, 2, 23]
- E. Campitelli. *metR: Tools for Easier Analysis of Meteorological Fields*, 2024. URL <https://CRAN.R-project.org/package=metR>. R package. [p27]
- G. Casella and R. L. Berger. *Statistical Inference*. Duxbury, 2nd edition, 2002. [p16]
- G. Csárdi. *cli: Helpers for Developing Command Line Interfaces*, 2024. URL <https://CRAN.R-project.org/package=cli>. R package. [p8]
- M. Greenwood. *A Report on the Natural Duration of Cancer*. Number 33 in Reports on Public Health and Medical Subjects. His Majesty's Stationery Office, London, 1926. [p1, 22]
- E. L. Kaplan and P. Meier. Nonparametric estimation from incomplete observations. *Journal of the American Statistical Association*, 53(282):457–481, 1958. doi: 10.1080/01621459.1958.10501452. [p1, 22]
- M. Kay. ggdist: Visualizations of distributions and uncertainty in the grammar of graphics. *IEEE Transactions on Visualization and Computer Graphics*, 30(1):983–993, 2024. doi: 10.1109/TVCG.2023.3327195. [p1, 27]
- P. Massart. The tight constant in the Dvoretzky–Kiefer–Wolfowitz inequality. *The Annals of Probability*, 18(3):1269–1283, 1990. doi: 10.1214/aop/1176990746. [p18, 25]
- J. R. Michael. The stabilized probability plot. *Biometrika*, 70(1):11–17, 1983. doi: 10.1093/biomet/70.1.11. URL <https://www.jstor.org/stable/2335939>. [p19]
- V. N. Nair. Confidence bands for survival functions with censored data: A comparative study. *Technometrics*, 26(3):265–275, 1984. doi: 10.1080/00401706.1984.10487964. [p1, 22]
- W. Nelson. Theory and applications of hazard plotting for censored failure data. *Technometrics*, 14(4): 945–966, 1972. doi: 10.1080/00401706.1972.10488991. [p1, 2, 23]

- J. Otto and D. Kahle. *ggdensity: Improved bivariate density visualization in R*. *The R Journal*, 15(2): 220–236, 2023. doi: 10.32614/RJ-2023-048. [p8, 9, 27]
- R. Pruim, D. T. Kaplan, and N. J. Horton. The mosaic package: Helping students to “think with data” using R. *The R Journal*, 9(1):77–102, 2017. doi: 10.32614/RJ-2017-024. [p1, 27]
- D. Sarkar. *Lattice: Multivariate Data Visualization with R*. Springer, New York, 2008. ISBN 978-0-387-75968-5. URL <https://lmdvr.r-forge.r-project.org>. [p1, 27]
- D. Turner, D. Kahle, and R. X. Sturdivant. *ggvfields: Vector Field Visualizations with ggplot2*, 2026. URL <https://github.com/dusty-turner/ggvfields>. R package version 0.1.0. [p6, 27]
- H. Wickham. A layered grammar of graphics. *Journal of Computational and Graphical Statistics*, 19(1): 3–28, 2010. doi: 10.1198/jcgs.2009.07098. [p1]
- H. Wickham. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York, 2nd edition, 2016. ISBN 978-3-319-24277-4. URL <https://ggplot2-book.org/>. [p1, 27]
- L. Wilkinson. *The Grammar of Graphics*. Springer-Verlag New York, 2nd edition, 2005. doi: 10.1007/0-387-28695-0. [p2]

Dusty Turner
United States Military Academy
West Point, NY
ORCID: 0000-0000-0000-0000
dusty.s.turner@gmail.com

David Kahle
Baylor University
Waco, TX
ORCID: 0000-0002-9999-1558
david@kahle.io

Rodney X. Sturdivant
Baylor University
Waco, TX
rodney_sturdivant@baylor.edu

James Otto
Alcon Inc.
Fort Worth, TX
james.otto@alcon.com